

Predicting the Sale Price of Bulldozers using Machine Learning

In this notebook, we're going to go through an example machine learning project with the goal of predicting the sale price of bulldozers.

Since we're trying to predict a number, this kind of problem is known as a **regression problem**.

The data and evaluation metric we'll be using (root mean square log error or RMSLE) is from the [Kaggle Bluebook for Bulldozers competition](#).

The techniques used in here have been inspired and adapted from the [fast.ai](#) machine learning course.

We'll work through each step and by the end of the notebook, we'll have a trained machine learning model which predicts the sale price of a bulldozer given different characteristics about it.

1. Problem Definition

For this dataset, the problem we're trying to solve, or better, the question we're trying to answer is,

How well can we predict the future sale price of a bulldozer, given its characteristics and previous examples of how much similar bulldozers have been sold for?

2. Data

Looking at the dataset from [Kaggle](#), it's a time series problem. This means there's a time attribute to dataset.

In this case, it's historical sales data of bulldozers. Including things like, model type, size, sale date and more.

There are 3 datasets:

- **Train.csv** - Historical bulldozer sales examples up to 2011 (close to 400,000 examples with 50+ different attributes, including SalePrice which is the **target variable**).
- **Valid.csv** - Historical bulldozer sales examples from January 1 2012 to April 30 2012 (close to 12,000 examples with the same attributes as **Train.csv**).
- **Test.csv** - Historical bulldozer sales examples from May 1 2012 to November 2012 (close to 12,000 examples but missing the SalePrice attribute, as this is what we'll be trying to predict).

3. Evaluation

For this problem, Kaggle has set the evaluation metric to being root mean squared log error ([RMSLE](#)) between the actual and predicted auction prices.

NOTE: As with many regression evaluations metric, the goal will be to get this value as low as possible (minimize the error).

For example, the goal for this project will be to build a machine learning model which minimizes RMSLE

To see how well the model is doing, we'll calculate the RMSLE and then compare the results to others on the [Kaggle leaderboard](#).

4. Features

Features are different parts of the data. During this step, we'll want to start finding out about the data.

One of the most common ways to do this, is to create a **data dictionary**.

For this dataset, Kaggle provide a [data dictionary](#) which contains information about what each attribute of the dataset means, this file is directly from the Kaggle competition page (account required) or view it on Google Sheets.

```
In [111]: # Import data analysis tools
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn
import seaborn as sns

# Model
from sklearn.ensemble import RandomForestRegressor
```

Read the Data

```
In [3]: # Import the training and validation set
df = pd.read_csv('data/bluebook-for-bulldozers/TrainAndValid.csv',
                 low_memory=False)
df.head()
```

```
Out[3]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	Usage
0	1139246	66000.0	999089	3157	121	3.0	2004	68.0	
1	1139248	57000.0	117657	77	121	3.0	1996	4640.0	
2	1139249	10000.0	434808	7009	121	3.0	2001	2838.0	
3	1139251	38500.0	1026470	332	121	3.0	2001	3486.0	
4	1139253	11000.0	1057373	17311	121	3.0	2007	722.0	M

5 rows × 53 columns

```
In [4]: # No parse_dates... check dtype of "saledate"
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 412698 entries, 0 to 412697
Data columns (total 53 columns):
```

#	Column	Non-Null Count	Dtype
0	SalesID	412698 non-null	int64
1	SalePrice	412698 non-null	float64
2	MachineID	412698 non-null	int64
3	ModelID	412698 non-null	int64
4	datasource	412698 non-null	int64
5	auctioneerID	392562 non-null	float64
6	YearMade	412698 non-null	int64
7	MachineHoursCurrentMeter	147504 non-null	float64
8	UsageBand	73670 non-null	object
9	saledate	412698 non-null	object
10	fiModelDesc	412698 non-null	object
11	fiBaseModel	412698 non-null	object
12	fiSecondaryDesc	271971 non-null	object
13	fiModelSeries	58667 non-null	object
14	fiModelDescriptor	74816 non-null	object
15	ProductSize	196093 non-null	object
16	fiProductClassDesc	412698 non-null	object
17	state	412698 non-null	object
18	ProductGroup	412698 non-null	object
19	ProductGroupDesc	412698 non-null	object
20	Drive_System	107087 non-null	object
21	Enclosure	412364 non-null	object
22	Forks	197715 non-null	object
23	Pad_Type	81096 non-null	object
24	Ride_Control	152728 non-null	object
25	Stick	81096 non-null	object
26	Transmission	188007 non-null	object
27	Turbocharged	81096 non-null	object
28	Blade_Extension	25983 non-null	object
29	Blade_Width	25983 non-null	object
30	Enclosure_Type	25983 non-null	object
31	Engine_Horsepower	25983 non-null	object
32	Hydraulics	330133 non-null	object
33	Pushblock	25983 non-null	object
34	Ripper	106945 non-null	object
35	Scarifier	25994 non-null	object
36	Tip_Control	25983 non-null	object
37	Tire_Size	97638 non-null	object
38	Coupler	220679 non-null	object
39	Coupler_System	44974 non-null	object
40	Grouser_Tracks	44875 non-null	object
41	Hydraulics_Flow	44875 non-null	object
42	Track_Type	102193 non-null	object
43	Undercarriage_Pad_Width	102916 non-null	object
44	Stick_Length	102261 non-null	object
45	Thumb	102332 non-null	object
46	Pattern_Changer	102261 non-null	object
47	Grouser_Type	102193 non-null	object
48	Backhoe_Mounting	80712 non-null	object
49	Blade_Type	81875 non-null	object
50	Travel_Controls	81877 non-null	object
51	Differential_Type	71564 non-null	object
52	Steering_Controls	71522 non-null	object

dtypes: float64(3), int64(5), object(45)

memory usage: 166.9+ MB

```
In [5]: # To check for missing null values
df.isna().sum()
```

```
Out[5]: SalesID      0
SalePrice    0
MachineID    0
ModelID      0
```

datasource	0
auctioneerID	20136
YearMade	0
MachineHoursCurrentMeter	265194
UsageBand	339028
saledate	0
fiModelDesc	0
fiBaseModel	0
fiSecondaryDesc	140727
fiModelSeries	354031
fiModelDescriptor	337882
ProductSize	216605
fiProductClassDesc	0
state	0
ProductGroup	0
ProductGroupDesc	0
Drive_System	305611
Enclosure	334
Forks	214983
Pad_Type	331602
Ride_Control	259970
Stick	331602
Transmission	224691
Turbocharged	331602
Blade_Extension	386715
Blade_Width	386715
Enclosure_Type	386715
Engine_Horsepower	386715
Hydraulics	82565
Pushblock	386715
Ripper	305753
Scarifier	386704
Tip_Control	386715
Tire_Size	315060
Coupler	192019
Coupler_System	367724
Grouser_Tracks	367823
Hydraulics_Flow	367823
Track_Type	310505
Undercarriage_Pad_Width	309782
Stick_Length	310437
Thumb	310366
Pattern_Changer	310437
Grouser_Type	310505
Backhoe_Mounting	331986
Blade_Type	330823
Travel_Controls	330821
Differential_Type	341134
Steering_Controls	341176

dtype: int64


In [6]: `df.columns`

Out[6]:

```
Index(['SalesID', 'SalePrice', 'MachineID', 'ModelID', 'datasource',
      'auctioneerID', 'YearMade', 'MachineHoursCurrentMeter', 'UsageBand',
      'saledate', 'fiModelDesc', 'fiBaseModel', 'fiSecondaryDesc',
      'fiModelSeries', 'fiModelDescriptor', 'ProductSize',
      'fiProductClassDesc', 'state', 'ProductGroup', 'ProductGroupDesc',
      'Drive_System', 'Enclosure', 'Forks', 'Pad_Type', 'Ride_Control',
      'Stick', 'Transmission', 'Turbocharged', 'Blade_Extension',
      'Blade_Width', 'Enclosure_Type', 'Engine_Horsepower', 'Hydraulics',
      'Pushblock', 'Ripper', 'Scarifier', 'Tip_Control', 'Tire_Size',
      'Coupler', 'Coupler_System', 'Grouser_Tracks', 'Hydraulics_Flow',
      'Track_Type', 'Undercarriage_Pad_Width', 'Stick_Length', 'Thumb',
      'Pattern_Changer', 'Grouser_Type', 'Backhoe_Mounting', 'Blade_Type',
```

```
'Travel_Controls', 'Differential_Type', 'Steering_Controls'],  
dtype='object')
```


```
In [7]: # Plot the first 1000 dataset  
fig, ax = plt.subplots()  
ax.scatter(df['saledate'][:1000], df['SalePrice'][:1000]); # (scatter plot: the x axis cor
```



```
In [8]: df.saledate[:1000]
```

```
Out[8]: 0      11/16/2006 0:00  
1       3/26/2004 0:00  
2       2/26/2004 0:00  
3       5/19/2011 0:00  
4       7/23/2009 0:00  
      ...  
995     7/16/2009 0:00  
996     6/14/2007 0:00  
997     9/22/2005 0:00  
998     7/28/2005 0:00  
999     6/16/2011 0:00  
Name: saledate, Length: 1000, dtype: object
```

```
In [9]: df.SalePrice.plot.hist();
```



Parsing dates

When working with time series data, we want to enrich the time & date component as much as possible.


```

23 Pad_Type 81096 non-null object
24 Ride_Control 152728 non-null object
25 Stick 81096 non-null object
26 Transmission 188007 non-null object
27 Turbocharged 81096 non-null object
28 Blade_Extension 25983 non-null object
29 Blade_Width 25983 non-null object
30 Enclosure_Type 25983 non-null object
31 Engine_Horsepower 25983 non-null object
32 Hydraulics 330133 non-null object
33 Pushblock 25983 non-null object
34 Ripper 106945 non-null object
35 Scarifier 25994 non-null object
36 Tip_Control 25983 non-null object
37 Tire_Size 97638 non-null object
38 Coupler 220679 non-null object
39 Coupler_System 44974 non-null object
40 Grouser_Tracks 44875 non-null object
41 Hydraulics_Flow 44875 non-null object
42 Track_Type 102193 non-null object
43 Undercarriage_Pad_Width 102916 non-null object
44 Stick_Length 102261 non-null object
45 Thumb 102332 non-null object
46 Pattern_Changer 102261 non-null object
47 Grouser_Type 102193 non-null object
48 Backhoe_Mounting 80712 non-null object
49 Blade_Type 81875 non-null object
50 Travel_Controls 81877 non-null object
51 Differential_Type 71564 non-null object
52 Steering_Controls 71522 non-null object
dtypes: datetime64[ns](1), float64(3), int64(5), object(44)
memory usage: 166.9+ MB

```

```

In [14]: # Plot the first 1000 dataset
fig, ax = plt.subplots()
ax.scatter(df['saledate'][:1000], df['SalePrice'][:1000]);

```



```

In [15]: df.head()

```

```

Out[15]:
   SalesID  SalePrice  MachineID  ModelID  datasource  auctioneerID  YearMade  MachineHoursCurrentMeter  Usage
0   1139246    66000.0    999089    3157         121           3.0      2004                    68.0
1   1139248    57000.0    117657     77         121           3.0      1996                   4640.0

```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand
2	1139249	10000.0	434808	7009	121	3.0	2001	2838.0	
3	1139251	38500.0	1026470	332	121	3.0	2001	3486.0	
4	1139253	11000.0	1057373	17311	121	3.0	2007	722.0	Medium

5 rows × 53 columns

In [16]:

```
df.head().T
```

Out[16]:

	0	1	2	3	4
SalesID	1139246	1139248	1139249	1139251	1139253
SalePrice	66000.0	57000.0	10000.0	38500.0	11000.0
MachineID	999089	117657	434808	1026470	1057373
ModelID	3157	77	7009	332	17311
datasource	121	121	121	121	121
auctioneerID	3.0	3.0	3.0	3.0	3.0
YearMade	2004	1996	2001	2001	2007
MachineHoursCurrentMeter	68.0	4640.0	2838.0	3486.0	722.0
UsageBand	Low	Low	High	High	Medium
saledate	2006-11-16 00:00:00	2004-03-26 00:00:00	2004-02-26 00:00:00	2011-05-19 00:00:00	2009-07-23 00:00:00
fiModelDesc	521D	950FII	226	PC120-6E	S175
fiBaseModel	521	950	226	PC120	S175
fiSecondaryDesc	D	F	NaN	NaN	NaN
fiModelSeries	NaN	II	NaN	-6E	NaN
fiModelDescriptor	NaN	NaN	NaN	NaN	NaN
ProductSize	NaN	Medium	NaN	Small	NaN
fiProductClassDesc	Wheel Loader - 110.0 to 120.0 Horsepower	Wheel Loader - 150.0 to 175.0 Horsepower	Skid Steer Loader - 1351.0 to 1601.0 Lb Operat...	Hydraulic Excavator, Track - 12.0 to 14.0 Metr...	Skid Steer Loader - 1601.0 to 1751.0 Lb Operat...
state	Alabama	North Carolina	New York	Texas	New York
ProductGroup	WL	WL	SSL	TEX	SSL
ProductGroupDesc	Wheel Loader	Wheel Loader	Skid Steer Loaders	Track Excavators	Skid Steer Loaders
Drive_System	NaN	NaN	NaN	NaN	NaN
Enclosure	EROPS w AC	EROPS w AC	OROPS	EROPS w AC	EROPS
Forks	None or Unspecified	None or Unspecified	None or Unspecified	NaN	None or Unspecified

	0	1	2	3	4
Pad_Type	NaN	NaN	NaN	NaN	NaN
Ride_Control	None or Unspecified	None or Unspecified	NaN	NaN	NaN
Stick	NaN	NaN	NaN	NaN	NaN
Transmission	NaN	NaN	NaN	NaN	NaN
Turbocharged	NaN	NaN	NaN	NaN	NaN
Blade_Extension	NaN	NaN	NaN	NaN	NaN
Blade_Width	NaN	NaN	NaN	NaN	NaN
Enclosure_Type	NaN	NaN	NaN	NaN	NaN
Engine_Horsepower	NaN	NaN	NaN	NaN	NaN
Hydraulics	2 Valve	2 Valve	Auxiliary	2 Valve	Auxiliary
Pushblock	NaN	NaN	NaN	NaN	NaN
Ripper	NaN	NaN	NaN	NaN	NaN
Scarifier	NaN	NaN	NaN	NaN	NaN
Tip_Control	NaN	NaN	NaN	NaN	NaN
Tire_Size	None or Unspecified	23.5	NaN	NaN	NaN
Coupler	None or Unspecified	None or Unspecified	None or Unspecified	None or Unspecified	None or Unspecified
Coupler_System	NaN	NaN	None or Unspecified	NaN	None or Unspecified
Grouser_Tracks	NaN	NaN	None or Unspecified	NaN	None or Unspecified
Hydraulics_Flow	NaN	NaN	Standard	NaN	Standard
Track_Type	NaN	NaN	NaN	NaN	NaN
Undercarriage_Pad_Width	NaN	NaN	NaN	NaN	NaN
Stick_Length	NaN	NaN	NaN	NaN	NaN
Thumb	NaN	NaN	NaN	NaN	NaN
Pattern_Changer	NaN	NaN	NaN	NaN	NaN
Grouser_Type	NaN	NaN	NaN	NaN	NaN
Backhoe_Mounting	NaN	NaN	NaN	NaN	NaN
Blade_Type	NaN	NaN	NaN	NaN	NaN
Travel_Controls	NaN	NaN	NaN	NaN	NaN
Differential_Type	Standard	Standard	NaN	NaN	NaN
Steering_Controls	Conventional	Conventional	NaN	NaN	NaN

In [17]: `df.saledate.head(20)`

Out[17]:

0	2006-11-16
1	2004-03-26

```
2      2004-02-26
3      2011-05-19
4      2009-07-23
5      2008-12-18
6      2004-08-26
7      2005-11-17
8      2009-08-27
9      2007-08-09
10     2008-08-21
11     2006-08-24
12     2005-10-20
13     2006-01-26
14     2006-01-03
15     2006-11-16
16     2007-06-14
17     2010-01-28
18     2006-03-09
19     2005-11-17
Name: saledate, dtype: datetime64[ns]
```

Sort DataFrame by saledate

As we're working on a time series problem and trying to predict future examples given past examples, it makes sense to sort our data by date (when working with time series, it's a good idea to sort it by date) .

```
In [18]: # Sort DataFrame in date order
df.sort_values(by=['saledate'], inplace=True, ascending=True)
df.saledate.head(20)
```

```
Out[18]: 205615    1989-01-17
274835    1989-01-31
141296    1989-01-31
212552    1989-01-31
62755     1989-01-31
54653     1989-01-31
81383     1989-01-31
204924    1989-01-31
135376    1989-01-31
113390    1989-01-31
113394    1989-01-31
116419    1989-01-31
32138     1989-01-31
127610    1989-01-31
76171     1989-01-31
127000    1989-01-31
128130    1989-01-31
127626    1989-01-31
55455     1989-01-31
55454     1989-01-31
Name: saledate, dtype: datetime64[ns]
```

```
In [19]: df.head(20)
```

```
Out[19]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter
205615	1646770	9500.0	1126363	8434	132	18.0	1974	NaN
274835	1821514	14000.0	1194089	10150	132	99.0	1980	NaN
141296	1505138	50000.0	1473654	4139	132	99.0	1978	NaN

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter
212552	1671174	16000.0	1327630	8591	132	99.0	1980	NaN
62755	1329056	22000.0	1336053	4089	132	99.0	1984	NaN
54653	1301884	23500.0	1182999	4123	132	99.0	1976	NaN
81383	1379228	31000.0	1082797	7620	132	99.0	1986	NaN
204924	1645390	11750.0	1527216	8202	132	99.0	1970	NaN
135376	1493279	63000.0	1363756	2759	132	99.0	1987	NaN
113390	1449549	13000.0	1289412	3356	132	99.0	1966	NaN
113394	1449555	10500.0	1102310	3356	132	99.0	1966	NaN
116419	1453775	20000.0	1514650	7008	132	99.0	1974	NaN
32138	1264985	20000.0	1204499	6788	132	99.0	1984	NaN
127610	1475641	23500.0	1194367	7277	132	99.0	1973	NaN
76171	1364654	14000.0	1270628	7289	132	99.0	1968	NaN
127000	1474844	11250.0	1279993	7257	132	99.0	1979	NaN
128130	1476264	29000.0	1245504	7277	132	99.0	1978	NaN
127626	1475662	22000.0	1242833	7277	132	99.0	1973	NaN
55455	1305337	17000.0	1517075	3356	132	99.0	1972	NaN
55454	1305336	17000.0	1236263	3356	132	99.0	1972	NaN

20 rows × 53 columns

Make a copy of the original DataFrame

Since we're going to be manipulating the data, we'll make a copy of the original DataFrame and perform our changes there.

This will keep the original DataFrame in tact if we need it again.

```
In [20]: # Make a copy of the original DataFrame to perform edits on
df_tmp = df.copy()
```

Add datetime parameters for saledate column

Why?

So we can enrich our dataset with as much information as possible.

Because we imported the data using read_csv() and we asked pandas to parse the dates using parse_dates=["saledate"], we can now access the different [datetime attributes](#) of the saledate column.

```
In [21]: # Add datetime parameters for saledate column
df_tmp['saleYear'] = df_tmp.saledate.dt.year
df_tmp['saleMonth'] = df_tmp.saledate.dt.month
df_tmp['saleDay'] = df_tmp.saledate.dt.day
df_tmp['saleDayofweek'] = df_tmp.saledate.dt.dayofweek
df_tmp['saleDayofyear'] = df_tmp.saledate.dt.dayofyear

# Drop original saledate(Now we've enriched our DataFrame with date time features, we can
df_tmp.drop('saledate', axis=1, inplace=True)
```

```
In [22]: df_tmp.head().T
```

Out[22]:		205615	274835	141296	212552	62755
	SalesID	1646770	1821514	1505138	1671174	1329056
	SalePrice	9500.0	14000.0	50000.0	16000.0	22000.0
	MachineID	1126363	1194089	1473654	1327630	1336053
	ModelID	8434	10150	4139	8591	4089
	datasource	132	132	132	132	132
	auctioneerID	18.0	99.0	99.0	99.0	99.0
	YearMade	1974	1980	1978	1980	1984
	MachineHoursCurrentMeter	NaN	NaN	NaN	NaN	NaN
	UsageBand	NaN	NaN	NaN	NaN	NaN
	fiModelDesc	TD20	A66	D7G	A62	D3B
	fiBaseModel	TD20	A66	D7	A62	D3
	fiSecondaryDesc	NaN	NaN	G	NaN	B
	fiModelSeries	NaN	NaN	NaN	NaN	NaN
	fiModelDescriptor	NaN	NaN	NaN	NaN	NaN
	ProductSize	Medium	NaN	Large	NaN	NaN
	fiProductClassDesc	Track Type Tractor, Dozer - 105.0 to 130.0 Horsepower	Wheel Loader - 120.0 to 135.0 Horsepower	Track Type Tractor, Dozer - 190.0 to 260.0 Horsepower	Wheel Loader - Unidentified	Track Type Tractor, Dozer - 20.0 to 75.0 Horsepower
	state	Texas	Florida	Florida	Florida	Florida
	ProductGroup	TTT	WL	TTT	WL	TTT
	ProductGroupDesc	Track Type Tractors	Wheel Loader	Track Type Tractors	Wheel Loader	Track Type Tractors

	205615	274835	141296	212552	62755
Drive_System	NaN	NaN	NaN	NaN	NaN
Enclosure	OROPS	OROPS	OROPS	EROPS	OROPS
Forks	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Pad_Type	NaN	NaN	NaN	NaN	NaN
Ride_Control	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Stick	NaN	NaN	NaN	NaN	NaN
Transmission	Direct Drive	NaN	Standard	NaN	Standard
Turbocharged	NaN	NaN	NaN	NaN	NaN
Blade_Extension	NaN	NaN	NaN	NaN	NaN
Blade_Width	NaN	NaN	NaN	NaN	NaN
Enclosure_Type	NaN	NaN	NaN	NaN	NaN
Engine_Horsepower	NaN	NaN	NaN	NaN	NaN
Hydraulics	2 Valve	2 Valve	2 Valve	2 Valve	2 Valve
Pushblock	NaN	NaN	NaN	NaN	NaN
Ripper	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Scarifier	NaN	NaN	NaN	NaN	NaN
Tip_Control	NaN	NaN	NaN	NaN	NaN
Tire_Size	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler_System	NaN	NaN	NaN	NaN	NaN
Grouser_Tracks	NaN	NaN	NaN	NaN	NaN
Hydraulics_Flow	NaN	NaN	NaN	NaN	NaN
Track_Type	NaN	NaN	NaN	NaN	NaN
Undercarriage_Pad_Width	NaN	NaN	NaN	NaN	NaN
Stick_Length	NaN	NaN	NaN	NaN	NaN
Thumb	NaN	NaN	NaN	NaN	NaN
Pattern_Changer	NaN	NaN	NaN	NaN	NaN
Grouser_Type	NaN	NaN	NaN	NaN	NaN
Backhoe_Mounting	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Blade_Type	Straight	NaN	Straight	NaN	PAT
Travel_Controls	None or Unspecified	NaN	None or Unspecified	NaN	Lever
Differential_Type	NaN	Standard	NaN	Standard	NaN

	205615	274835	141296	212552	62755
Steering_Controls	NaN	Conventional	NaN	Conventional	NaN
saleYear	1989	1989	1989	1989	1989
saleMonth	1	1	1	1	1
saleDay	17	31	31	31	31
saleDayofweek	1	1	1	1	1
saleDayofyear	17	31	31	31	31

In [23]:

```
# Check the different values of different columns
df_tmp.state.value_counts()
```

Out[23]:

```
Florida      67320
Texas        53110
California   29761
Washington   16222
Georgia      14633
Maryland     13322
Mississippi  13240
Ohio         12369
Illinois     11540
Colorado     11529
New Jersey   11156
North Carolina 10636
Tennessee    10298
Alabama      10292
Pennsylvania 10234
South Carolina 9951
Arizona      9364
New York     8639
Connecticut  8276
Minnesota    7885
Missouri     7178
Nevada       6932
Louisiana    6627
Kentucky     5351
Maine        5096
Indiana      4124
Arkansas     3933
New Mexico   3631
Utah         3046
Unspecified  2801
Wisconsin    2745
New Hampshire 2738
Virginia     2353
Idaho        2025
Oregon       1911
Michigan     1831
Wyoming      1672
Iowa         1336
Montana      1336
Oklahoma     1326
Nebraska     866
West Virginia 840
Kansas       667
Delaware     510
North Dakota 480
Alaska       430
Massachusetts 347
Vermont      300
```

```
South Dakota      244
Hawaii            118
Rhode Island      83
Puerto Rico      42
Washington DC     2
Name: state, dtype: int64
```

5. Modelling

We've explored our dataset a little as well as enriched it with some datetime attributes, now let's try to model.

Why model so early?

We know the evaluation metric we're heading towards. We could spend more time doing exploratory data analysis (EDA), finding more out about the data ourselves but what we'll do instead is use a machine learning model to help us do EDA.

Following the [Scikit-Learn machine learning map](#), we find a `RandomForestRegressor()` might be a good candidate.

Note: Remember, one of the biggest goals of starting any new machine learning project is reducing the time between experiments.

```
In [24]: # Check for missing categories and different datatypes
df_tmp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 412698 entries, 205615 to 409203
Data columns (total 57 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   SalesID                              412698 non-null  int64
1   SalePrice                            412698 non-null  float64
2   MachineID                            412698 non-null  int64
3   ModelID                              412698 non-null  int64
4   datasource                           412698 non-null  int64
5   auctioneerID                         392562 non-null  float64
6   YearMade                             412698 non-null  int64
7   MachineHoursCurrentMeter             147504 non-null  float64
8   UsageBand                            73670 non-null   object
9   fiModelDesc                          412698 non-null  object
10  fiBaseModel                          412698 non-null  object
11  fiSecondaryDesc                      271971 non-null  object
12  fiModelSeries                       58667 non-null   object
13  fiModelDescriptor                    74816 non-null   object
14  ProductSize                         196093 non-null  object
15  fiProductClassDesc                  412698 non-null  object
16  state                               412698 non-null  object
17  ProductGroup                        412698 non-null  object
18  ProductGroupDesc                    412698 non-null  object
19  Drive_System                        107087 non-null  object
20  Enclosure                           412364 non-null  object
21  Forks                               197715 non-null  object
22  Pad_Type                            81096 non-null   object
23  Ride_Control                        152728 non-null  object
24  Stick                               81096 non-null   object
25  Transmission                        188007 non-null  object
26  Turbocharged                       81096 non-null   object
27  Blade_Extension                     25983 non-null   object
28  Blade_Width                         25983 non-null   object
29  Enclosure_Type                     25983 non-null   object
```

30	Engine_Horsepower	25983	non-null	object
31	Hydraulics	330133	non-null	object
32	Pushblock	25983	non-null	object
33	Ripper	106945	non-null	object
34	Scarifier	25994	non-null	object
35	Tip_Control	25983	non-null	object
36	Tire_Size	97638	non-null	object
37	Coupler	220679	non-null	object
38	Coupler_System	44974	non-null	object
39	Grouser_Tracks	44875	non-null	object
40	Hydraulics_Flow	44875	non-null	object
41	Track_Type	102193	non-null	object
42	Undercarriage_Pad_Width	102916	non-null	object
43	Stick_Length	102261	non-null	object
44	Thumb	102332	non-null	object
45	Pattern_Changer	102261	non-null	object
46	Grouser_Type	102193	non-null	object
47	Backhoe_Mounting	80712	non-null	object
48	Blade_Type	81875	non-null	object
49	Travel_Controls	81877	non-null	object
50	Differential_Type	71564	non-null	object
51	Steering_Controls	71522	non-null	object
52	saleYear	412698	non-null	int64
53	saleMonth	412698	non-null	int64
54	saleDay	412698	non-null	int64
55	saleDayofweek	412698	non-null	int64
56	saleDayofyear	412698	non-null	int64

dtypes: float64(3), int64(10), object(44)

memory usage: 182.6+ MB

```
In [25]: # Check for missing values
df_tmp.isna().sum()
```

```
Out[25]: SalesID          0
SalePrice          0
MachineID          0
ModelID            0
datasource         0
auctioneerID       20136
YearMade           0
MachineHoursCurrentMeter  265194
UsageBand          339028
fiModelDesc        0
fiBaseModel         0
fiSecondaryDesc     140727
fiModelSeries       354031
fiModelDescriptor   337882
ProductSize         216605
fiProductClassDesc  0
state              0
ProductGroup        0
ProductGroupDesc    0
Drive_System        305611
Enclosure           334
Forks               214983
Pad_Type            331602
Ride_Control        259970
Stick               331602
Transmission        224691
Turbocharged        331602
Blade_Extension     386715
Blade_Width         386715
Enclosure_Type      386715
Engine_Horsepower   386715
Hydraulics          82565
```



```

Pushblock      386715
Ripper         305753
Scarifier      386704
Tip_Control    386715
Tire_Size      315060
Coupler        192019
Coupler_System 367724
Grouser_Tracks 367823
Hydraulics_Flow 367823
Track_Type     310505
Undercarriage_Pad_Width 309782
Stick_Length   310437
Thumb          310366
Pattern_Changer 310437
Grouser_Type   310505
Backhoe_Mounting 331986
Blade_Type     330823
Travel_Controls 330821
Differential_Type 341134
Steering_Controls 341176
saleYear       0
saleMonth      0
saleDay        0
saleDayofweek  0
saleDayofyear  0
dtype: int64

```

Convert strings to categories

One way to help turn all of our data into numbers is to convert the columns with the string datatype into a category datatype.

To do this we can use the [pandas types API](#) which allows us to interact and manipulate the types of data.

```
In [26]: df_tmp.head().T
```

```
Out[26]:
```

	205615	274835	141296	212552	62755
SalesID	1646770	1821514	1505138	1671174	1329056
SalePrice	9500.0	14000.0	50000.0	16000.0	22000.0
MachineID	1126363	1194089	1473654	1327630	1336053
ModelID	8434	10150	4139	8591	4089
datasource	132	132	132	132	132
auctioneerID	18.0	99.0	99.0	99.0	99.0
YearMade	1974	1980	1978	1980	1984
MachineHoursCurrentMeter	NaN	NaN	NaN	NaN	NaN
UsageBand	NaN	NaN	NaN	NaN	NaN
fiModelDesc	TD20	A66	D7G	A62	D3B
fiBaseModel	TD20	A66	D7	A62	D3
fiSecondaryDesc	NaN	NaN	G	NaN	B
fiModelSeries	NaN	NaN	NaN	NaN	NaN
fiModelDescriptor	NaN	NaN	NaN	NaN	NaN

	205615	274835	141296	212552	62755
ProductSize	Medium	NaN	Large	NaN	NaN
fiProductClassDesc	Track Type Tractor, Dozer - 105.0 to 130.0 Horsepower	Wheel Loader - 120.0 to 135.0 Horsepower	Track Type Tractor, Dozer - 190.0 to 260.0 Horsepower	Wheel Loader - Unidentified	Track Type Tractor, Dozer - 20.0 to 75.0 Horsepower
state	Texas	Florida	Florida	Florida	Florida
ProductGroup	TTT	WL	TTT	WL	TTT
ProductGroupDesc	Track Type Tractors	Wheel Loader	Track Type Tractors	Wheel Loader	Track Type Tractors
Drive_System	NaN	NaN	NaN	NaN	NaN
Enclosure	OROPS	OROPS	OROPS	EROPS	OROPS
Forks	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Pad_Type	NaN	NaN	NaN	NaN	NaN
Ride_Control	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Stick	NaN	NaN	NaN	NaN	NaN
Transmission	Direct Drive	NaN	Standard	NaN	Standard
Turbocharged	NaN	NaN	NaN	NaN	NaN
Blade_Extension	NaN	NaN	NaN	NaN	NaN
Blade_Width	NaN	NaN	NaN	NaN	NaN
Enclosure_Type	NaN	NaN	NaN	NaN	NaN
Engine_Horsepower	NaN	NaN	NaN	NaN	NaN
Hydraulics	2 Valve	2 Valve	2 Valve	2 Valve	2 Valve
Pushblock	NaN	NaN	NaN	NaN	NaN
Ripper	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Scarifier	NaN	NaN	NaN	NaN	NaN
Tip_Control	NaN	NaN	NaN	NaN	NaN
Tire_Size	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler_System	NaN	NaN	NaN	NaN	NaN
Grouser_Tracks	NaN	NaN	NaN	NaN	NaN
Hydraulics_Flow	NaN	NaN	NaN	NaN	NaN
Track_Type	NaN	NaN	NaN	NaN	NaN
Undercarriage_Pad_Width	NaN	NaN	NaN	NaN	NaN
Stick_Length	NaN	NaN	NaN	NaN	NaN
Thumb	NaN	NaN	NaN	NaN	NaN

	205615	274835	141296	212552	62755
Pattern_Changer	NaN	NaN	NaN	NaN	NaN
Grouser_Type	NaN	NaN	NaN	NaN	NaN
Backhoe_Mounting	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Blade_Type	Straight	NaN	Straight	NaN	PAT
Travel_Controls	None or Unspecified	NaN	None or Unspecified	NaN	Lever
Differential_Type	NaN	Standard	NaN	Standard	NaN
Steering_Controls	NaN	Conventional	NaN	Conventional	NaN
saleYear	1989	1989	1989	1989	1989
saleMonth	1	1	1	1	1
saleDay	17	31	31	31	31
saleDayofweek	1	1	1	1	1
saleDayofyear	17	31	31	31	31

In [27]: `pd.api.types.is_string_dtype(df_tmp['UsageBand'])`

Out[27]: True

In [28]:

```
# Find the columns which contain strings
for label, content in df_tmp.items():
    if pd.api.types.is_string_dtype(content):
        print(label)
```

```
UsageBand
fiModelDesc
fiBaseModel
fiSecondaryDesc
fiModelSeries
fiModelDescriptor
ProductSize
fiProductClassDesc
state
ProductGroup
ProductGroupDesc
Drive_System
Enclosure
Forks
Pad_Type
Ride_Control
Stick
Transmission
Turbocharged
Blade_Extension
Blade_Width
Enclosure_Type
Engine_Horsepower
Hydraulics
Pushblock
Ripper
Scarifier
Tip_Control
```



```

34 Scarifier                25994 non-null    category
35 Tip_Control              25983 non-null    category
36 Tire_Size               97638 non-null    category
37 Coupler                 220679 non-null    category
38 Coupler_System          44974 non-null    category
39 Grouser_Tracks          44875 non-null    category
40 Hydraulics_Flow         44875 non-null    category
41 Track_Type              102193 non-null    category
42 Undercarriage_Pad_Width 102916 non-null    category
43 Stick_Length            102261 non-null    category
44 Thumb                   102332 non-null    category
45 Pattern_Changer         102261 non-null    category
46 Grouser_Type            102193 non-null    category
47 Backhoe_Mounting        80712 non-null     category
48 Blade_Type              81875 non-null     category
49 Travel_Controls         81877 non-null     category
50 Differential_Type        71564 non-null     category
51 Steering_Controls        71522 non-null     category
52 saleYear                412698 non-null    int64
53 saleMonth               412698 non-null    int64
54 saleDay                 412698 non-null    int64
55 saleDayofweek           412698 non-null    int64
56 saleDayofyear           412698 non-null    int64
dtypes: category(44), float64(3), int64(10)
memory usage: 63.2 MB

```

```
In [32]: df_tmp.state.cat.categories
```

```

Out[32]: Index(['Alabama', 'Alaska', 'Arizona', 'Arkansas', 'California', 'Colorado',
      'Connecticut', 'Delaware', 'Florida', 'Georgia', 'Hawaii', 'Idaho',
      'Illinois', 'Indiana', 'Iowa', 'Kansas', 'Kentucky', 'Louisiana',
      'Maine', 'Maryland', 'Massachusetts', 'Michigan', 'Minnesota',
      'Mississippi', 'Missouri', 'Montana', 'Nebraska', 'Nevada',
      'New Hampshire', 'New Jersey', 'New Mexico', 'New York',
      'North Carolina', 'North Dakota', 'Ohio', 'Oklahoma', 'Oregon',
      'Pennsylvania', 'Puerto Rico', 'Rhode Island', 'South Carolina',
      'South Dakota', 'Tennessee', 'Texas', 'Unspecified', 'Utah', 'Vermont',
      'Virginia', 'Washington', 'Washington DC', 'West Virginia', 'Wisconsin',
      'Wyoming'],
      dtype='object')

```

```
In [33]: df_tmp.state.cat.codes
```

```

Out[33]: 205615    43
274835     8
141296     8
212552     8
62755      8
..
410879     4
412476     4
411927     4
407124     4
409203     4
Length: 412698, dtype: int8

```

All of our data is categorical and thus we can now turn the categories into numbers, however it's still missing values...

```
In [34]: # Check missing data
df_tmp.isnull().sum()/len(df_tmp)
```

```
Out[34]: SalesID                0.000000
```

SalePrice	0.000000
MachineID	0.000000
ModelID	0.000000
datasource	0.000000
auctioneerID	0.048791
YearMade	0.000000
MachineHoursCurrentMeter	0.642586
UsageBand	0.821492
fiModelDesc	0.000000
fiBaseModel	0.000000
fiSecondaryDesc	0.340993
fiModelSeries	0.857845
fiModelDescriptor	0.818715
ProductSize	0.524851
fiProductClassDesc	0.000000
state	0.000000
ProductGroup	0.000000
ProductGroupDesc	0.000000
Drive_System	0.740520
Enclosure	0.000809
Forks	0.520921
Pad_Type	0.803498
Ride_Control	0.629928
Stick	0.803498
Transmission	0.544444
Turbocharged	0.803498
Blade_Extension	0.937041
Blade_Width	0.937041
Enclosure_Type	0.937041
Engine_Horsepower	0.937041
Hydraulics	0.200062
Pushblock	0.937041
Ripper	0.740864
Scarifier	0.937014
Tip_Control	0.937041
Tire_Size	0.763415
Coupler	0.465277
Coupler_System	0.891024
Grouser_Tracks	0.891264
Hydraulics_Flow	0.891264
Track_Type	0.752378
Undercarriage_Pad_Width	0.750626
Stick_Length	0.752213
Thumb	0.752041
Pattern_Changer	0.752213
Grouser_Type	0.752378
Backhoe_Mounting	0.804428
Blade_Type	0.801610
Travel_Controls	0.801606
Differential_Type	0.826595
Steering_Controls	0.826697
saleYear	0.000000
saleMonth	0.000000
saleDay	0.000000
saleDayofweek	0.000000
saleDayofyear	0.000000
dtype: float64	

In the format it's in, it's still good to be worked with, let's save it to file and reimport it so we can continue on.

Save Processed Data

```
In [35]: # Save preprocessed data
df_tmp.to_csv('data/bluebook-for-bulldozers/train_tmp.csv', index=False)
```

In [36]:

```
# Import preprocessed data
df_tmp = pd.read_csv('data/bluebook-for-bulldozers/train_tmp.csv', low_memory=False)
df_tmp.head().T
```

Out[36]:

	0	1	2	3	4
SalesID	1646770	1821514	1505138	1671174	1329056
SalePrice	9500.0	14000.0	50000.0	16000.0	22000.0
MachineID	1126363	1194089	1473654	1327630	1336053
ModelID	8434	10150	4139	8591	4089
datasource	132	132	132	132	132
auctioneerID	18.0	99.0	99.0	99.0	99.0
YearMade	1974	1980	1978	1980	1984
MachineHoursCurrentMeter	NaN	NaN	NaN	NaN	NaN
UsageBand	NaN	NaN	NaN	NaN	NaN
fiModelDesc	TD20	A66	D7G	A62	D3B
fiBaseModel	TD20	A66	D7	A62	D3
fiSecondaryDesc	NaN	NaN	G	NaN	B
fiModelSeries	NaN	NaN	NaN	NaN	NaN
fiModelDescriptor	NaN	NaN	NaN	NaN	NaN
ProductSize	Medium	NaN	Large	NaN	NaN
fiProductClassDesc	Track Type Tractor, Dozer - 105.0 to 130.0 Horsepower	Wheel Loader - 120.0 to 135.0 Horsepower	Track Type Tractor, Dozer - 190.0 to 260.0 Horsepower	Wheel Loader - Unidentified	Track Type Tractor, Dozer - 20.0 to 75.0 Horsepower
state	Texas	Florida	Florida	Florida	Florida
ProductGroup	TTT	WL	TTT	WL	TTT
ProductGroupDesc	Track Type Tractors	Wheel Loader	Track Type Tractors	Wheel Loader	Track Type Tractors
Drive_System	NaN	NaN	NaN	NaN	NaN
Enclosure	OROPS	OROPS	OROPS	EROPS	OROPS
Forks	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Pad_Type	NaN	NaN	NaN	NaN	NaN
Ride_Control	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Stick	NaN	NaN	NaN	NaN	NaN
Transmission	Direct Drive	NaN	Standard	NaN	Standard
Turbocharged	NaN	NaN	NaN	NaN	NaN
Blade_Extension	NaN	NaN	NaN	NaN	NaN
Blade_Width	NaN	NaN	NaN	NaN	NaN
Enclosure_Type	NaN	NaN	NaN	NaN	NaN

	0	1	2	3	4
Engine_Horsepower	NaN	NaN	NaN	NaN	NaN
Hydraulics	2 Valve	2 Valve	2 Valve	2 Valve	2 Valve
Pushblock	NaN	NaN	NaN	NaN	NaN
Ripper	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Scarifier	NaN	NaN	NaN	NaN	NaN
Tip_Control	NaN	NaN	NaN	NaN	NaN
Tire_Size	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler_System	NaN	NaN	NaN	NaN	NaN
Grouser_Tracks	NaN	NaN	NaN	NaN	NaN
Hydraulics_Flow	NaN	NaN	NaN	NaN	NaN
Track_Type	NaN	NaN	NaN	NaN	NaN
Undercarriage_Pad_Width	NaN	NaN	NaN	NaN	NaN
Stick_Length	NaN	NaN	NaN	NaN	NaN
Thumb	NaN	NaN	NaN	NaN	NaN
Pattern_Changer	NaN	NaN	NaN	NaN	NaN
Grouser_Type	NaN	NaN	NaN	NaN	NaN
Backhoe_Mounting	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Blade_Type	Straight	NaN	Straight	NaN	PAT
Travel_Controls	None or Unspecified	NaN	None or Unspecified	NaN	Lever
Differential_Type	NaN	Standard	NaN	Standard	NaN
Steering_Controls	NaN	Conventional	NaN	Conventional	NaN
saleYear	1989	1989	1989	1989	1989
saleMonth	1	1	1	1	1
saleDay	17	31	31	31	31
saleDayofweek	1	1	1	1	1
saleDayofyear	17	31	31	31	31

Excellent, our processed DataFrame has the columns we added to it but it's still missing values.

```
In [37]: # Check for missing values
df_tmp.isna().sum()
```

```
Out[37]: SalesID          0
SalePrice          0
MachineID          0
ModelID            0
```


datasource	0
auctioneerID	20136
YearMade	0
MachineHoursCurrentMeter	265194
UsageBand	339028
fiModelDesc	0
fiBaseModel	0
fiSecondaryDesc	140727
fiModelSeries	354031
fiModelDescriptor	337882
ProductSize	216605
fiProductClassDesc	0
state	0
ProductGroup	0
ProductGroupDesc	0
Drive_System	305611
Enclosure	334
Forks	214983
Pad_Type	331602
Ride_Control	259970
Stick	331602
Transmission	224691
Turbocharged	331602
Blade_Extension	386715
Blade_Width	386715
Enclosure_Type	386715
Engine_Horsepower	386715
Hydraulics	82565
Pushblock	386715
Ripper	305753
Scarifier	386704
Tip_Control	386715
Tire_Size	315060
Coupler	192019
Coupler_System	367724
Grouser_Tracks	367823
Hydraulics_Flow	367823
Track_Type	310505
Undercarriage_Pad_Width	309782
Stick_Length	310437
Thumb	310366
Pattern_Changer	310437
Grouser_Type	310505
Backhoe_Mounting	331986
Blade_Type	330823
Travel_Controls	330821
Differential_Type	341134
Steering_Controls	341176
saleYear	0
saleMonth	0
saleDay	0
saleDayofweek	0
saleDayofyear	0

dtype: int64

Fill missing values

From experience with machine learning models. We know two things:

1. All of our data has to be numerical
2. There can't be any missing values

And as we've seen using `df_tmp.isna().sum()` our data still has plenty of missing values.

Let's fill them.

Filling numerical values first

We're going to fill any column with missing values with the median of that column.

```
In [40]: # Check for the columns that have numeric values
for label, content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        print(label)
```

```
SalesID
SalePrice
MachineID
ModelID
datasource
auctioneerID
YearMade
MachineHoursCurrentMeter
saleYear
saleMonth
saleDay
saleDayofweek
saleDayofyear
```

```
In [41]: # Check for which numeric columns have null(missing) values
for label, content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)
```

```
auctioneerID
MachineHoursCurrentMeter
```

```
In [42]: # Fill numeric rows with the median
for label, content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():

            # Add a binary column which tells us if the data was missing or not
            df_tmp[label+'_is_missing'] = pd.isnull(content)

            # Fill missing numeric values with median since it's more robust than the mean
            df_tmp[label] = content.fillna(content.median())
```

Why add a binary column indicating whether the data was missing or not?

We can easily fill all of the missing numeric values in our dataset with the median. However, a numeric value may be missing for a reason. In other words, absence of evidence may be evidence of absence. Adding a binary column which indicates whether the value was missing or not helps to retain this information.

```
In [43]: # Check if there's any null values
for label, content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)
```

```
In [44]: # Check to see how many examples were missing
```

```
df_tmp.auctioneerID_is_missing.value_counts()
```

```
Out[44]: False      392562
         True       20136
         Name: auctioneerID_is_missing, dtype: int64
```

```
In [45]: df_tmp.isna().sum()
```

```
Out[45]: SalesID      0
         SalePrice    0
         MachineID    0
         ModelID      0
         datasource    0
         auctioneerID  0
         YearMade      0
         MachineHoursCurrentMeter  0
         UsageBand    339028
         fiModelDesc    0
         fiBaseModel    0
         fiSecondaryDesc  140727
         fiModelSeries  354031
         fiModelDescriptor  337882
         ProductSize    216605
         fiProductClassDesc  0
         state         0
         ProductGroup    0
         ProductGroupDesc  0
         Drive_System    305611
         Enclosure      334
         Forks          214983
         Pad_Type       331602
         Ride_Control    259970
         Stick          331602
         Transmission    224691
         Turbocharged    331602
         Blade_Extension  386715
         Blade_Width     386715
         Enclosure_Type  386715
         Engine_Horsepower  386715
         Hydraulics      82565
         Pushblock       386715
         Ripper          305753
         Scarifier       386704
         Tip_Control     386715
         Tire_Size       315060
         Coupler         192019
         Coupler_System  367724
         Grouser_Tracks  367823
         Hydraulics_Flow  367823
         Track_Type      310505
         Undercarriage_Pad_Width  309782
         Stick_Length    310437
         Thumb          310366
         Pattern_Changer  310437
         Grouser_Type     310505
         Backhoe_Mounting  331986
         Blade_Type      330823
         Travel_Controls  330821
         Differential_Type  341134
         Steering_Controls  341176
         saleYear        0
         saleMonth       0
         saleDay         0
         saleDayofweek    0
         saleDayofyear    0
```

```
auctioneerID_is_missing      0
MachineHoursCurrentMeter_is_missing  0
dtype: int64
```

Filling and turning categorical variables to numbers

Now we've filled the numeric values, we'll do the same with the categorical values at the same time as turning them into numbers.

In [46]:

```
# Check columns which *aren't* numeric
for label, content in df_tmp.items():
    if not pd.api.types.is_numeric_dtype(content):
        print(label)
```

```
UsageBand
fiModelDesc
fiBaseModel
fiSecondaryDesc
fiModelSeries
fiModelDescriptor
ProductSize
fiProductClassDesc
state
ProductGroup
ProductGroupDesc
Drive_System
Enclosure
Forks
Pad_Type
Ride_Control
Stick
Transmission
Turbocharged
Blade_Extension
Blade_Width
Enclosure_Type
Engine_Horsepower
Hydraulics
Pushblock
Ripper
Scarifier
Tip_Control
Tire_Size
Coupler
Coupler_System
Grouser_Tracks
Hydraulics_Flow
Track_Type
Undercarriage_Pad_Width
Stick_Length
Thumb
Pattern_Changer
Grouser_Type
Backhoe_Mounting
Blade_Type
Travel_Controls
Differential_Type
Steering_Controls
```

In [47]:

```
# Turn categorical variables into numbers
for label, content in df_tmp.items():
    # Check columns which *aren't* numeric
    if not pd.api.types.is_numeric_dtype(content):
```

```
# Add binary column to indicate whether sample had missing value
df_tmp[label + '_is_missing'] = pd.isnull(content)

# Turn categories into numbers and add +1 because pandas encodes missing categories as -1
df_tmp[label] = pd.Categorical(content).codes + 1
```

In [48]: `df_tmp.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 412698 entries, 0 to 412697
Columns: 103 entries, SalesID to Steering_Controls_is_missing
dtypes: bool(46), float64(3), int16(4), int64(10), int8(40)
memory usage: 77.9 MB
```

In [49]: `df_tmp.isna().sum()`

```
Out[49]: SalesID          0
SalePrice          0
MachineID          0
ModelID            0
datasource         0
...
Backhoe_Mounting_is_missing  0
Blade_Type_is_missing      0
Travel_Controls_is_missing  0
Differential_Type_is_missing 0
Steering_Controls_is_missing 0
Length: 103, dtype: int64
```

In [50]: `df_tmp.head().T`

```
Out[50]:
```

	0	1	2	3	4
SalesID	1646770	1821514	1505138	1671174	1329056
SalePrice	9500.0	14000.0	50000.0	16000.0	22000.0
MachineID	1126363	1194089	1473654	1327630	1336053
ModelID	8434	10150	4139	8591	4089
datasource	132	132	132	132	132
...
Backhoe_Mounting_is_missing	False	True	False	True	False
Blade_Type_is_missing	False	True	False	True	False
Travel_Controls_is_missing	False	True	False	True	False
Differential_Type_is_missing	True	False	True	False	True
Steering_Controls_is_missing	True	False	True	False	True

103 rows × 5 columns

Now all of our data is numeric and there are no missing values, we should be able to build a machine learning model!

Let's reinstantiate [RandomForestRegressor](#).

This will take a few minutes which is too long for interacting with it. So what we'll do is create a subset of rows to work with.

```
In [51]: len(df_tmp)
```

```
Out[51]: 412698
```

```
In [56]: %%time

# Instantiate model
model = RandomForestRegressor(n_jobs=-1,
                             random_state=42)

# Fit the model
model.fit(df_tmp.drop('SalePrice', axis = 1), df_tmp.SalePrice)
```

Wall time: 8min 47s

```
Out[56]: RandomForestRegressor(n_jobs=-1, random_state=42)
```

```
In [57]: # Score the model
model.score(df_tmp.drop('SalePrice', axis = 1), df_tmp.SalePrice)
```

```
Out[57]: 0.9875468079970562
```

The metric above is not reliable.

Splitting data into train/validation sets

```
In [58]: df_tmp.head()
```

```
Out[58]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	Usage
0	1646770	9500.0	1126363	8434	132	18.0	1974		0.0
1	1821514	14000.0	1194089	10150	132	99.0	1980		0.0
2	1505138	50000.0	1473654	4139	132	99.0	1978		0.0
3	1671174	16000.0	1327630	8591	132	99.0	1980		0.0
4	1329056	22000.0	1336053	4089	132	99.0	1984		0.0

5 rows × 103 columns

According to the [Kaggle data page](#), the validation set and test set are split according to dates.

This makes sense since we're working on a time series problem.

E.g. using past events to try and predict future events.

Knowing this, randomly splitting our data into train and test sets using something like `train_test_split()` wouldn't work.

Instead, we split our data into training, validation and test sets using the date each sample occurred.

In our case:

- Training = all samples up until 2011
- Valid = all samples from January 1, 2012 - April 30, 2012
- Test = all samples from May 1, 2012 - November 2012

In [59]: `df_tmp.saleYear`

Out[59]:

0	1989
1	1989
2	1989
3	1989
4	1989
	...
412693	2012
412694	2012
412695	2012
412696	2012
412697	2012

Name: saleYear, Length: 412698, dtype: int64

In [60]: `df_tmp.saleYear.value_counts()`

Out[60]:

2009	43849
2008	39767
2011	35197
2010	33390
2007	32208
2006	21685
2005	20463
2004	19879
2001	17594
2000	17415
2002	17246
2003	15254
1998	13046
1999	12793
2012	11573
1997	9785
1996	8829
1995	8530
1994	7929
1993	6303
1992	5519
1991	5109
1989	4806
1990	4529

Name: saleYear, dtype: int64

In [61]:

```
# Split the data into training and validation set
df_val = df_tmp[df_tmp.saleYear == 2012]
df_train = df_tmp[df_tmp.saleYear != 2012]

len(df_val), len(df_train)
```

Out[61]: (11573, 401125)

In [62]:

```
# Split the data into X and y
X_train, y_train = df_train.drop('SalePrice', axis = 1), df_train.SalePrice
X_valid, y_valid = df_val.drop('SalePrice', axis = 1), df_val.SalePrice
```

```
X_train.shape, y_train.shape, X_valid.shape, y_valid.shape
```

```
Out[62]: ((401125, 102), (401125,), (11573, 102), (11573,))
```

```
In [63]: y_train
```

```
Out[63]: 0          9500.0
1         14000.0
2         50000.0
3         16000.0
4         22000.0
...
401120     29000.0
401121     11000.0
401122     11000.0
401123     18000.0
401124     13500.0
Name: SalePrice, Length: 401125, dtype: float64
```

Building an evaluation function

According to Kaggle for the Bluebook for Bulldozers competition, the [evaluation function](#) they use is root mean squared log error (RMSLE).

RMSLE = generally you don't care as much if you're off by \$10 as much as you'd care if you were off by 10%, you care more about ratios rather than differences. **MAE** (mean absolute error) is more about exact differences.

It's important to understand the evaluation metric you're going for.

Since Scikit-Learn doesn't have a function built-in for RMSLE, we'll create our own.

We can do this by taking the square root of Scikit-Learn's [mean_squared_log_error](#) (MSLE). MSLE is the same as taking the log of mean squared error (MSE).

We'll also calculate the MAE and R^2 .

```
In [67]: # Create evaluation function (the competition uses Root Mean Square Log Error RMSLE)
from sklearn.metrics import mean_squared_log_error, mean_absolute_error

def rmsle(y_test, y_preds):
    """
    calculates root mean squared log error between predictions
    and true labels
    """
    return np.sqrt(mean_squared_log_error(y_test, y_preds))

# Create a function to evaluate our model
def show_scores(model):
    train_preds = model.predict(X_train)
    val_preds = model.predict(X_valid)
    scores = {'Training MAE': mean_absolute_error(y_train, train_preds),
              'Valid MAE': mean_absolute_error(y_valid, val_preds),
              'Training RMSLE': rmsle(y_train, train_preds),
              'Valid RMSLE': rmsle(y_valid, val_preds),
              'Training R^2': model.score(X_train, y_train),
              'Valid R^2': model.score(X_valid, y_valid)}
    return scores
```


Testing our model on a subset (to tune the hyperparameters)

Retraining an entire model would take far too long to continuing experimenting as fast as we want to.

So what we'll do is take a sample of the training set and tune the hyperparameters on that before training a larger model.

If your experiments are taking longer than 10-seconds (give or take how long you have to wait), you should be trying to speed things up. You can speed things up by sampling less data or using a faster computer.

```
In [65]: # This takes far too long for experimenting...

# %%time
# # Retrain a model on training data
# model.fit(X_train, y_train)
# show_scores(model)
```

```
In [70]: len(X_train)
```

```
Out[70]: 401125
```

Depending on your computer, making calculations on ~400,000 rows may take a while...

Let's alter the number of samples each `n_estimator` in the `RandomForestRegressor` see's using the `max_samples` parameter.

```
In [71]: # Change max samples in RandomForestRegressor
model = RandomForestRegressor(n_jobs=-1,
                             random_state=42,
                             max_samples=10000)
```

Setting `max_samples` to 10000 means every `n_estimator` (default 100) in our `RandomForestRegressor` will only see 10000 random samples from our `DataFrame` instead of the entire 400,000.

In other words, we'll be looking at 40x less samples which means we'll get faster computation speeds but we should expect our results to worsen (simple the model has less samples to learn patterns from).

```
In [72]: %%time
# Cutting down the max number of samples each tree can see improves training time
model.fit(X_train, y_train)
```

Wall time: 25.3 s

```
Out[72]: RandomForestRegressor(max_samples=10000, n_jobs=-1, random_state=42)
```

```
In [73]: show_scores(model)
```

```
Out[73]: {'Training MAE': 5561.2988092240585,
          'Valid MAE': 7177.26365505919,
          'Training RMSLE': 0.257745378256977,
          'Valid RMSLE': 0.29362638671089003,
          'Training R^2': 0.8606658995199189,
          'Valid R^2': 0.8320374995090507}
```

Hyperparameter tuning with RandomizedSearchCV

You can increase `n_iter` to try more combinations of hyperparameters but in our case, we'll try 20 and see where it gets us.

In [74]:

```
%%time
from sklearn.model_selection import RandomizedSearchCV

# Different RandomForestRegressor hyperparameters
rf_grid = {'n_estimators': np.arange(10, 100, 10),
           'max_depth': [None, 3, 5, 10],
           'min_samples_split': np.arange(2, 20, 2),
           'min_samples_leaf': np.arange(1, 20, 2),
           'max_features': [0.5, 1, 'sqrt', 'auto'],
           'max_samples': [10000]}

# Instantiate RandomizedSearchCV model
rs_model = RandomizedSearchCV(RandomForestRegressor(),
                              param_distributions = rf_grid,
                              n_iter=20,
                              cv=5,
                              verbose=True)

# Fit the RandomizedSearchCV model
rs_model.fit(X_train, y_train)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits
Wall time: 12min

Out[74]:

```
RandomizedSearchCV(cv=5, estimator=RandomForestRegressor(), n_iter=20,
                  param_distributions={'max_depth': [None, 3, 5, 10],
                                      'max_features': [0.5, 1, 'sqrt',
                                                      'auto'],
                                      'max_samples': [10000],
                                      'min_samples_leaf': array([ 1,  3,  5,  7,  9, 11,
                                                                13, 15, 17, 19]),
                                      'min_samples_split': array([ 2,  4,  6,  8, 10, 1
                                                                2, 14, 16, 18]),
                                      'n_estimators': array([10, 20, 30, 40, 50, 60, 70,
                                                                80, 90])},
                  verbose=True)
```

In [75]:

```
# Find the best parameters from the RandomizedSearch
rs_model.best_params_
```

Out[75]:

```
{'n_estimators': 80,
 'min_samples_split': 16,
 'min_samples_leaf': 7,
 'max_samples': 10000,
 'max_features': 0.5,
 'max_depth': None}
```

In [76]:

```
# Evaluate the RandomizedSearch model
show_scores(rs_model)
```

Out[76]:

```
{'Training MAE': 6220.152642892111,
 'Valid MAE': 7523.711301067172,
 'Training RMSLE': 0.2816262592419515,
 'Valid RMSLE': 0.30412742636021506,
 'Training R^2': 0.8276368097596425,
 'Valid R^2': 0.8153446700929028}
```

Train a model with the best parameters

In [80]:

```
%%time
```

```
# Most ideal hyperparameters
ideal_model = RandomForestRegressor(n_estimators=90,
                                   min_samples_leaf=1,
                                   min_samples_split=14,
                                   max_features=0.5,
                                   n_jobs=-1,
                                   max_samples=None,
                                   random_state = 42)

# Fit the ideal model
ideal_model.fit(X_train, y_train)
```

Wall time: 3min 45s

Out[80]: RandomForestRegressor(max_features=0.5, min_samples_split=14, n_estimators=90, n_jobs=-1, random_state=42)

In [81]: *# scores for ideal model (trained on all the data)*
show_scores(ideal_model)

Out[81]: {'Training MAE': 2929.921568433354,
'Valid MAE': 5926.828732131214,
'Training RMSLE': 0.14357234549234102,
'Valid RMSLE': 0.24477262723855048,
'Training R^2': 0.9596218418812977,
'Valid R^2': 0.8829608391304742}

In [82]: *# scores for rs_model (only trained on ~10,000 samples)*
show_scores(rs_model)

Out[82]: {'Training MAE': 6220.152642892111,
'Valid MAE': 7523.711301067172,
'Training RMSLE': 0.2816262592419515,
'Valid RMSLE': 0.30412742636021506,
'Training R^2': 0.8276368097596425,
'Valid R^2': 0.8153446700929028}

Make predictions on test data

Now we've got a trained model, it's time to make predictions on the test data.

Remember what we've done.

Our model is trained on data prior to 2011. However, the test data is from May 1 2012 to November 2012.

So what we're doing is trying to use the patterns our model has learned in the training data to predict the sale price of a Bulldozer with characteristics it's never seen before but are assumed to be similar to that of those in the training data.

In [90]: *# Import the test data*
df_test = pd.read_csv('data/bluebook-for-bulldozers/Test.csv',
 parse_dates=['saledate'])

df_test

Out[90]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand
0	1227829	1006309	3168	121	3	1999	3688.0	Low
1	1227844	1022817	7271	121	3	1000	28555.0	High

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand
2	1227847	1031560	22805	121	3	2004	6038.0	Medium
3	1227848	56204	1269	121	3	2006	8940.0	High
4	1227863	1053887	22312	121	3	2005	2286.0	Low
...
12452	6643171	2558317	21450	149	2	2008	NaN	NaN
12453	6643173	2558332	21434	149	2	2005	NaN	NaN
12454	6643184	2558342	21437	149	2	1000	NaN	NaN
12455	6643186	2558343	21437	149	2	2006	NaN	NaN
12456	6643196	2558346	21446	149	2	2008	NaN	NaN

12457 rows × 52 columns

In [91]:

```
df_test.head()
```

Out[91]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	sale
0	1227829	1006309	3168	121	3	1999	3688.0	Low	2000
1	1227844	1022817	7271	121	3	1000	28555.0	High	2000
2	1227847	1031560	22805	121	3	2004	6038.0	Medium	2000
3	1227848	56204	1269	121	3	2006	8940.0	High	2000
4	1227863	1053887	22312	121	3	2005	2286.0	Low	2000

5 rows × 52 columns

Preprocessing the data(getting the test data in the same format as our training data)

Our model has been trained on data formatted in the same way as the training data.

This means in order to make predictions on the test data, we need to take the same steps we used to preprocess the training data to preprocess the test data.

Remember: Whatever you do to the training data, you have to do to the test data.

Let's create a function for doing so (by copying the preprocessing steps we used above).

```
In [92]: def preprocess_data(df):
        """
        performs transformations on df and returns transformed df
        """
        # Add datetime parameters for saledate
        df["saleYear"] = df.saledate.dt.year
        df["saleMonth"] = df.saledate.dt.month
        df["saleDay"] = df.saledate.dt.day
        df["saleDayofweek"] = df.saledate.dt.dayofweek
        df["saleDayofyear"] = df.saledate.dt.dayofyear

        # Drop original saledate
        df.drop("saledate", axis=1, inplace=True)

        # Fill numeric rows with the median
        for label, content in df.items():
            if pd.api.types.is_numeric_dtype(content):
                if pd.isnull(content).sum():

                    # Add a binary column which tells us if the data was missing or not
                    df[label+"_is_missing"] = pd.isnull(content)

                    # Fill missing numeric values with median since it's more robust than the
                    df[label] = content.fillna(content.median())

            # Turn categorical missing variables into numbers
            if not pd.api.types.is_numeric_dtype(content):
                df[label+"_is_missing"] = pd.isnull(content)

                # We add the +1 because pandas encodes missing categories as -1
                df[label] = pd.Categorical(content).codes+1

        return df
```

Now we've got a function for preprocessing data, let's preprocess the test dataset into the same format as our training dataset.

```
In [93]: # Process the test data
df_test = preprocess_data(df_test)
df_test.head()
```

```
Out[93]:
```

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiM
0	1227829	1006309	3168	121	3	1999	3688.0	2	
1	1227844	1022817	7271	121	3	1000	28555.0	1	
2	1227847	1031560	22805	121	3	2004	6038.0	3	
3	1227848	56204	1269	121	3	2006	8940.0	1	
4	1227863	1053887	22312	121	3	2005	2286.0	2	

5 rows × 101 columns

```
In [94]: x_train.head()
```

```
Out[94]:
```

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiM
0	1646770	1126363	8434	132	18.0	1974	0.0	0	
1	1821514	1194089	10150	132	99.0	1980	0.0	0	

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiM
2	1505138	1473654	4139	132	99.0	1978	0.0	0	
3	1671174	1327630	8591	132	99.0	1980	0.0	0	
4	1329056	1336053	4089	132	99.0	1984	0.0	0	

5 rows × 102 columns

```
In [95]: # We can find how the columns differ using sets
set(X_train.columns) - set(df_test.columns)
```

```
Out[95]: {'auctioneerID_is_missing'}
```

In this case, it's because the test dataset wasn't missing any auctioneerID fields.

To fix it, we'll add a column to the test dataset called auctioneerID_is_missing and fill it with False, since none of the auctioneerID fields are missing in the test dataset.

```
In [96]: # Match test dataset columns to training dataset
df_test['auctioneerID_is_missing'] = False
df_test
```

```
Out[96]:
```

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiM
0	1227829	1006309	3168	121	3	1999	3688.0	2	
1	1227844	1022817	7271	121	3	1000	28555.0	1	
2	1227847	1031560	22805	121	3	2004	6038.0	3	
3	1227848	56204	1269	121	3	2006	8940.0	1	
4	1227863	1053887	22312	121	3	2005	2286.0	2	
...
12452	6643171	2558317	21450	149	2	2008	3525.0	0	
12453	6643173	2558332	21434	149	2	2005	3525.0	0	
12454	6643184	2558342	21437	149	2	1000	3525.0	0	
12455	6643186	2558343	21437	149	2	2006	3525.0	0	
12456	6643196	2558346	21446	149	2	2008	3525.0	0	

12457 rows × 102 columns

Now the test dataset matches the training dataset, we should be able to make predictions on it using our trained model.

```
In [97]: # Make predictions on the test dataset using the best model
test_preds = ideal_model.predict(df_test)
```

When looking at the [Kaggle submission requirements](#), we see that if we wanted to make a submission, the data is required to be in a certain format. Namely, a DataFrame containing the SalesID and the predicted SalePrice of the bulldozer.

```
In [98]:
```

```
# Create DataFrame compatible with Kaggle submission requirements
df_preds = pd.DataFrame()
df_preds['SalesID'] = df_test['SalesID']
df_preds['SalesPrice'] = test_preds
df_preds
```

```
Out[98]:
```

	SalesID	SalesPrice
0	1227829	19789.714023
1	1227844	19565.299439
2	1227847	48800.694515
3	1227848	61935.439690
4	1227863	42767.772640
...
12452	6643171	42265.421622
12453	6643173	14782.962987
12454	6643184	15604.288740
12455	6643186	19025.395893
12456	6643196	29334.359121

12457 rows × 2 columns

```
In [99]: # Export to CSV
df_preds.to_csv('data/bluebook-for-bulldozers/predictions.csv', index = False)
```

Feature Importance

Feature importance seeks to figure out which different attributes of the data were most important when it comes to predicting the target variable.

```
In [113]: # Find feature importance of our best model
ideal_model.feature_importances_
```

```
Out[113]: array([3.36536104e-02, 2.06759219e-02, 4.12493898e-02, 1.65309161e-03,
        3.37161218e-03, 2.06931736e-01, 3.23666071e-03, 1.03610744e-03,
        4.53494594e-02, 4.41122835e-02, 6.01968492e-02, 4.71102363e-03,
        1.59697733e-02, 1.50896732e-01, 4.57018149e-02, 5.95545126e-03,
        2.22480465e-03, 3.33266898e-03, 2.95703302e-03, 6.73951018e-02,
        6.65158539e-04, 3.98094940e-05, 9.84704968e-04, 2.85986069e-04,
        1.12813990e-03, 2.47446180e-05, 9.54863612e-04, 8.81712425e-03,
        1.92386650e-03, 2.12645204e-03, 3.62305282e-03, 3.44945894e-03,
        3.27317013e-03, 1.77639716e-03, 1.14087894e-03, 5.80942535e-03,
        8.12450169e-04, 1.31354263e-02, 1.39777375e-03, 3.28177035e-03,
        1.28211228e-03, 8.69687238e-04, 2.05251738e-03, 5.97963087e-04,
        5.41678527e-04, 3.54317048e-04, 3.88821556e-04, 2.41962294e-03,
        8.59586252e-04, 2.92501513e-04, 2.15302259e-04, 7.33725591e-02,
        3.79418026e-03, 5.68832526e-03, 2.90982570e-03, 9.82376367e-03,
        2.59268029e-04, 1.74795886e-03, 3.23057317e-04, 0.00000000e+00,
        0.00000000e+00, 2.70655122e-03, 1.26563042e-03, 3.88600447e-03,
        3.21688076e-02, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
        0.00000000e+00, 1.02259045e-04, 1.08902429e-05, 3.10665050e-04,
        3.99741175e-06, 1.06257032e-04, 3.58165848e-06, 3.05858606e-04,
        1.69755228e-05, 1.39801322e-03, 1.86677952e-03, 3.62449442e-03,
```

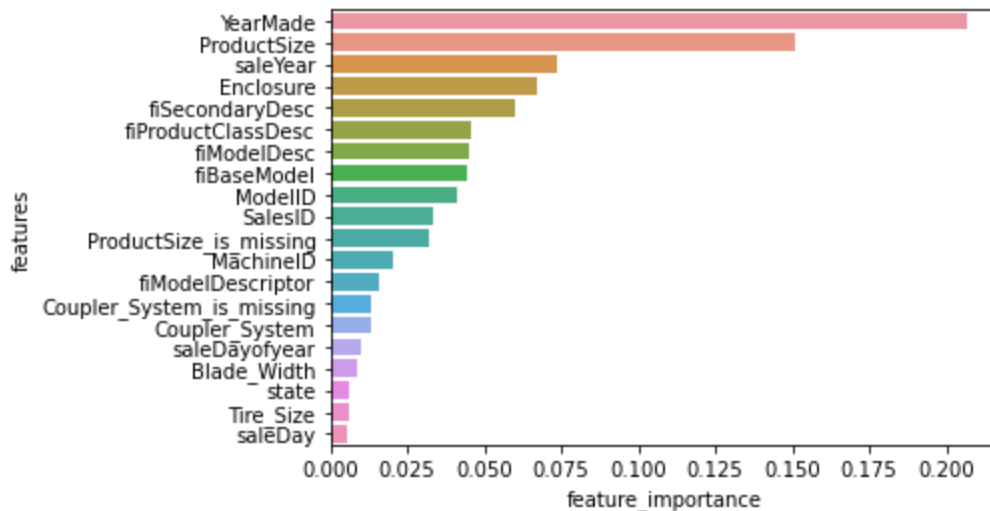
```
2.26585416e-04, 3.61935932e-03, 9.31853135e-04, 1.33423840e-03,
1.34288116e-03, 1.12092114e-03, 2.71203804e-03, 2.71750836e-04,
1.34054763e-02, 1.76499457e-03, 1.11672051e-03, 4.27793962e-05,
1.84939697e-04, 4.09157804e-05, 8.90131799e-05, 4.49277953e-05,
4.21445103e-05, 2.97559133e-04, 1.86343248e-04, 1.60013050e-04,
1.02015171e-04, 1.30937225e-04])
```

In [115...

```
# Helper function for plotting feature importance
def plot_features(columns, importances, n=20):
    df = (pd.DataFrame({'features': columns,
                        'feature_importance': importances})
          .sort_values('feature_importance', ascending=False)
          .reset_index(drop=True))
    sns.barplot(x='feature_importance',
                y='features',
                data=df[:n],
                orient='h')
```

In [116...

```
plot_features(X_train.columns, ideal_model.feature_importances_)
```



In [117...

```
sum(ideal_model.feature_importances_)
```

Out[117...

1.0

In [105...

```
df.ProductSize.isna().sum()
```

Out[105...

216605

In [106...

```
df.ProductSize.value_counts()
```

Out[106...

```
Medium          64342
Large / Medium  51297
Small           27057
Mini            25721
Large            21396
Compact          6280
Name: ProductSize, dtype: int64
```

In [107...

```
df.Turbocharged.value_counts()
```

Out[107...

```
None or Unspecified    77111
```



```
Yes 3985
Name: Turbocharged, dtype: int64
```

```
In [108... df.Thumb.value_counts()
```

```
Out[108... None or Unspecified    85074
Manual                9678
Hydraulic             7580
Name: Thumb, dtype: int64
```